

# IMT 573: lab-dplyr

Alexander Davis

Tuesday, October 13, 2020

## Obectives

In this demo we will practice working with data. We will employ the `dplyr` verbs to manipulate a dataset in various ways.

```
# Load some helpful libraries for this course
library(tidyverse)
```

**Import and inspect the data** We'll be using Gapminder data, which represents the health and wealth of nations. It was pioneered by Hans Rosling, who is famous for describing the prosperity of nations over time through famines, wars and other historic events with this beautiful data visualization in his 2006 TED Talk: The best stats you've ever seen:

Let's import this data into R and see what it looks like.

```
gapminder <- read_csv('https://raw.githubusercontent.com/OHI-Science/data-science-training/master/data/')

##
## -- Column specification -----
## cols(
##   country = col_character(),
##   year = col_double(),
##   pop = col_double(),
##   continent = col_character(),
##   lifeExp = col_double(),
##   gdpPercap = col_double()
## )
```

What is the size of this dataset?

```
# Find and print the number of row and columns in this dataset
dim(gapminder)
```

```
## [1] 1704    6
```

Consider looking at the raw data. What variables are listed? What data types are used for this data? Do you spot any immediate concerns?

```
# Use RStudio utils to view the raw dataset
# View(gapminder)
# Get variable summaries
summary(gapminder)
```

```
##      country      year      pop      continent
## Length:1704      Min.   :1952      Min.   :6.001e+04      Length:1704
## Class :character  1st Qu.:1966      1st Qu.:2.794e+06      Class :character
## Mode  :character  Median :1980      Median :7.024e+06      Mode  :character
##                               Mean   :1980      Mean   :2.960e+07
##                               3rd Qu.:1993      3rd Qu.:1.959e+07
##                               Max.   :2007      Max.   :1.319e+09
##      lifeExp      gdpPercap
## Min.   :23.60      Min.   : 241.2
## 1st Qu.:48.20      1st Qu.: 1202.1
## Median :60.71      Median : 3531.8
## Mean   :59.47      Mean   : 7215.3
## 3rd Qu.:70.85      3rd Qu.: 9325.5
## Max.   :82.60      Max.   :113523.1
```

## dplyr Verbs for Data Manipulation

dplyr is a grammar of data manipulation, providing a consistent set of verbs that help you solve the most common data manipulation challenges:

- `mutate()` adds new variables that are functions of existing variables
- `select()` picks variables based on their names.
- `filter()` picks cases based on their values.
- `summarise()` reduces multiple values down to a single summary.
- `arrange()` changes the ordering of the rows.

Let's practice our data manipulation skills with the gapminder data that you just loaded.

### Filter: keep rows matching criteria

Q1: Filter the gapminder data for life expectancy less than 29

```
filter(gapminder, lifeExp < 29)
```

```
## # A tibble: 2 x 6
##   country      year      pop continent lifeExp gdpPercap
##   <chr>      <dbl>   <dbl> <chr>      <dbl>    <dbl>
## 1 Afghanistan 1952 8425333 Asia      28.8     779.
## 2 Rwanda      1992 7290203 Africa    23.6     737.
```

Q2: "Filter the gapminder data for the country Mexico"

```
## Your turn
filter(gapminder, country == "Mexico")
```

```
## # A tibble: 12 x 6
##   country year      pop continent lifeExp gdpPercap
##   <chr>   <dbl>    <dbl> <chr>      <dbl>    <dbl>
## 1 Mexico  1952  30144317 Americas    50.8     3478.
## 2 Mexico  1957  35015548 Americas    55.2     4132.
## 3 Mexico  1962  41121485 Americas    58.3     4582.
## 4 Mexico  1967  47995559 Americas    60.1     5755.
## 5 Mexico  1972  55984294 Americas    62.4     6809.
## 6 Mexico  1977  63759976 Americas    65.0     7675.
## 7 Mexico  1982  71640904 Americas    67.4     9611.
## 8 Mexico  1987  80122492 Americas    69.5     8688.
## 9 Mexico  1992  88111030 Americas    71.5     9472.
## 10 Mexico 1997  95895146 Americas    73.7     9767.
## 11 Mexico 2002 102479927 Americas    74.9    10742.
## 12 Mexico 2007 108700891 Americas    76.2    11978.
```

Q3: if we want two country names? We can't use the == operator here, because it can only operate on one thing at a time. We will use the %in% operator:

```
filter(gapminder, country %in% c("Mexico", "Peru"))
```

```
## # A tibble: 24 x 6
##   country year      pop continent lifeExp gdpPercap
##   <chr>   <dbl>    <dbl> <chr>      <dbl>    <dbl>
## 1 Mexico  1952  30144317 Americas    50.8     3478.
## 2 Mexico  1957  35015548 Americas    55.2     4132.
## 3 Mexico  1962  41121485 Americas    58.3     4582.
## 4 Mexico  1967  47995559 Americas    60.1     5755.
## 5 Mexico  1972  55984294 Americas    62.4     6809.
## 6 Mexico  1977  63759976 Americas    65.0     7675.
## 7 Mexico  1982  71640904 Americas    67.4     9611.
## 8 Mexico  1987  80122492 Americas    69.5     8688.
## 9 Mexico  1992  88111030 Americas    71.5     9472.
## 10 Mexico 1997  95895146 Americas    73.7     9767.
## # ... with 14 more rows
```

Q4: "We want data for Mexico in 2002?"

```
## Your turn
filter(gapminder, country == "Mexico" & year == 2002)
```

```
## # A tibble: 1 x 6
##   country year      pop continent lifeExp gdpPercap
##   <chr>   <dbl>    <dbl> <chr>      <dbl>    <dbl>
## 1 Mexico  2002 102479927 Americas    74.9    10742.
```

Select: pick columns by name

```
select(gapminder, year, country, lifeExp)
```

```
## # A tibble: 1,704 x 3
##   year country    lifeExp
##   <dbl> <chr>      <dbl>
## 1  1952 Afghanistan  28.8
## 2  1957 Afghanistan  30.3
## 3  1962 Afghanistan  32.0
## 4  1967 Afghanistan  34.0
## 5  1972 Afghanistan  36.1
## 6  1977 Afghanistan  38.4
## 7  1982 Afghanistan  39.9
## 8  1987 Afghanistan  40.8
## 9  1992 Afghanistan  41.7
## 10 1997 Afghanistan  41.8
## # ... with 1,694 more rows
```

We can also use - to deselect columns

```
select(gapminder, -continent, -lifeExp) # you can use - to deselect columns
```

```
## # A tibble: 1,704 x 4
##   country    year    pop gdpPercap
##   <chr>      <dbl>  <dbl>    <dbl>
## 1 Afghanistan 1952 8425333    779.
## 2 Afghanistan 1957 9240934    821.
## 3 Afghanistan 1962 10267083    853.
## 4 Afghanistan 1967 11537966    836.
## 5 Afghanistan 1972 13079460    740.
## 6 Afghanistan 1977 14880372    786.
## 7 Afghanistan 1982 12881816    978.
## 8 Afghanistan 1987 13867957    852.
## 9 Afghanistan 1992 16317921    649.
## 10 Afghanistan 1997 22227415    635.
## # ... with 1,694 more rows
```

**Arrange:** reorder rows

Q: Sorted by year and then life-expectancy

```
arrange(gapminder, year, lifeExp)
```

```
## # A tibble: 1,704 x 6
##   country    year    pop continent lifeExp gdpPercap
##   <chr>      <dbl>  <dbl> <chr>      <dbl>    <dbl>
## 1 Afghanistan 1952 8425333 Asia      28.8     779.
## 2 Gambia      1952 284320 Africa    30       485.
## 3 Angola      1952 4232095 Africa   30.0    3521.
## 4 Sierra Leone 1952 2143249 Africa   30.3     880.
## 5 Mozambique   1952 6446316 Africa   31.3     469.
## 6 Burkina Faso 1952 4469979 Africa   32.0     543.
## 7 Guinea-Bissau 1952 580653 Africa   32.5     300.
## 8 Yemen Rep.   1952 4963829 Asia     32.5     782.
## 9 Somalia      1952 2526994 Africa   33.0    1136.
```

```
## 10 Guinea          1952 2664249 Africa      33.6      510.
## # ... with 1,694 more rows
```

Q: But your boss wants to see the data sorted in reverse chronological order.

```
## Your turn
arrange(gapminder, -year)
```

```
## # A tibble: 1,704 x 6
##   country      year      pop continent lifeExp gdpPercap
##   <chr>      <dbl>    <dbl> <chr>      <dbl>    <dbl>
## 1 Afghanistan  2007  31889923 Asia      43.8     975.
## 2 Albania      2007   3600523 Europe    76.4   5937.
## 3 Algeria      2007  33333216 Africa    72.3   6223.
## 4 Angola       2007  12420476 Africa    42.7   4797.
## 5 Argentina    2007  40301927 Americas  75.3  12779.
## 6 Australia    2007  20434176 Oceania   81.2  34435.
## 7 Austria      2007   8199783 Europe    79.8   36126.
## 8 Bahrain      2007    708573 Asia     75.6   29796.
## 9 Bangladesh   2007  150448339 Asia     64.1    1391.
## 10 Belgium     2007   10392226 Europe    79.4   33693.
## # ... with 1,694 more rows
```

**Mutate: add new variables**

Q: Imagine we want to know each country's annual GDP. We can multiply pop by gdpPercap to create a new column named gdp.

```
gapminder %>%
  mutate(gdp = pop * gdpPercap)
```

```
## # A tibble: 1,704 x 7
##   country      year      pop continent lifeExp gdpPercap      gdp
##   <chr>      <dbl>    <dbl> <chr>      <dbl>    <dbl>    <dbl>
## 1 Afghanistan  1952  8425333 Asia      28.8     779.  6567086330.
## 2 Afghanistan  1957  9240934 Asia      30.3     821.  7585448670.
## 3 Afghanistan  1962 10267083 Asia      32.0     853.  8758855797.
## 4 Afghanistan  1967 11537966 Asia      34.0     836.  9648014150.
## 5 Afghanistan  1972 13079460 Asia      36.1     740.  9678553274.
## 6 Afghanistan  1977 14880372 Asia      38.4     786. 11697659231.
## 7 Afghanistan  1982 12881816 Asia      39.9     978. 12598563401.
## 8 Afghanistan  1987 13867957 Asia      40.8     852. 11820990309.
## 9 Afghanistan  1992 16317921 Asia      41.7     649. 10595901589.
## 10 Afghanistan 1997 22227415 Asia      41.8     635. 14121995875.
## # ... with 1,694 more rows
```

Q. Now we want to calculate the annual GDP for all Asian countries in the year 2007 and add it as a new column. How can you do it?

```
## Your turn
gapminder %>%
  filter(continent == "Asia" & year == 2007) %>%
  mutate(gdp = pop * gdpPercap)
```

```
## # A tibble: 33 x 7
##   country      year      pop continent lifeExp gdpPercap      gdp
##   <chr>      <dbl>    <dbl> <chr>      <dbl>    <dbl>    <dbl>
## 1 Afghanistan  2007  31889923 Asia      43.8      975.  3.11e10
## 2 Bahrain      2007   708573 Asia      75.6    29796.  2.11e10
## 3 Bangladesh   2007 150448339 Asia      64.1    1391.  2.09e11
## 4 Cambodia     2007  14131858 Asia      59.7    1714.  2.42e10
## 5 China        2007 1318683096 Asia      73.0    4959.  6.54e12
## 6 Hong Kong China 2007   6980412 Asia      82.2   39725.  2.77e11
## 7 India        2007 1110396331 Asia      64.7    2452.  2.72e12
## 8 Indonesia    2007  223547000 Asia      70.6    3541.  7.92e11
## 9 Iran         2007   69453570 Asia      71.0   11606.  8.06e11
## 10 Iraq        2007   27499638 Asia      59.5    4471.  1.23e11
## # ... with 23 more rows
```

Q. Now we want to calculate the population in thousands for all Asian countries in the year 2007 and add it as a new column. How can you do it? *Hint: You will use the same logic as the previous question, just with gdp calculation replaced with pop/1000 calculation*

```
## Your turn
gapminder %>%
  filter(continent == "Asia" & year == 2007) %>%
  mutate(popThousands = pop/1000)
```

```
## # A tibble: 33 x 7
##   country      year      pop continent lifeExp gdpPercap popThousands
##   <chr>      <dbl>    <dbl> <chr>      <dbl>    <dbl>    <dbl>
## 1 Afghanistan  2007  31889923 Asia      43.8      975.    31890.
## 2 Bahrain      2007   708573 Asia      75.6    29796.    709.
## 3 Bangladesh   2007 150448339 Asia      64.1    1391.   150448.
## 4 Cambodia     2007  14131858 Asia      59.7    1714.   14132.
## 5 China        2007 1318683096 Asia      73.0    4959.  1318683.
## 6 Hong Kong China 2007   6980412 Asia      82.2   39725.    6980.
## 7 India        2007 1110396331 Asia      64.7    2452.  1110396.
## 8 Indonesia    2007  223547000 Asia      70.6    3541.   223547.
## 9 Iran         2007   69453570 Asia      71.0   11606.   69454.
## 10 Iraq        2007   27499638 Asia      59.5    4471.   27500.
## # ... with 23 more rows
```

## Summarize with group\_by

Q. Find the total population on each continent in 2005

```
gapminder %>%
  filter(year == 2002) %>%
  group_by(continent) %>%
  mutate(cont_pop = sum(pop))
```

```
## # A tibble: 142 x 7
## # Groups:   continent [5]
##   country      year      pop continent lifeExp gdpPercap  cont_pop
##   <chr>      <dbl>    <dbl> <chr>      <dbl>    <dbl>    <dbl>
```

```
## 1 Afghanistan 2002 25268405 Asia 42.1 727. 3601802203
## 2 Albania 2002 3508512 Europe 75.7 4604. 578223869
## 3 Algeria 2002 31287142 Africa 71.0 5288. 833723916
## 4 Angola 2002 10866106 Africa 41.0 2773. 833723916
## 5 Argentina 2002 38331121 Americas 74.3 8798. 849772762
## 6 Australia 2002 19546792 Oceania 80.4 30688. 23454829
## 7 Austria 2002 8148312 Europe 79.0 32418. 578223869
## 8 Bahrain 2002 656397 Asia 74.8 23404. 3601802203
## 9 Bangladesh 2002 135656790 Asia 62.0 1136. 3601802203
## 10 Belgium 2002 10311970 Europe 78.3 30486. 578223869
## # ... with 132 more rows
```

Q. Find the median population on each continent in 2002

*## Your Turn*

```
gapminder %>%
  select(continent, country, year, pop) %>%
  filter(year == 2002) %>%
  group_by(continent, year) %>%
  summarise(median = median(pop))
```

```
## 'summarise()' regrouping output by 'continent' (override with '.groups' argument)
```

```
## # A tibble: 5 x 3
## # Groups:   continent [5]
##   continent year median
##   <chr>    <dbl>   <dbl>
## 1 Africa    2002  8821778.
## 2 Americas  2002  8650322
## 3 Asia      2002 22662365
## 4 Europe    2002  9518744
## 5 Oceania   2002 11727414.
```

## Summarize with group\_by

We can use more than one grouping variable. Let's get total populations by continent and year.

```
gapminder %>%
  group_by(continent, year) %>%
  summarize(cont_pop = sum(pop))
```

```
## 'summarise()' regrouping output by 'continent' (override with '.groups' argument)
```

```
## # A tibble: 60 x 3
## # Groups:   continent [5]
##   continent year cont_pop
##   <chr>    <dbl>   <dbl>
## 1 Africa    1952 237640501
## 2 Africa    1957 264837738
## 3 Africa    1962 296516865
## 4 Africa    1967 335289489
```

```
## 5 Africa      1972 379879541
## 6 Africa      1977 433061021
## 7 Africa      1982 499348587
## 8 Africa      1987 574834110
## 9 Africa      1992 659081517
## 10 Africa     1997 743832984
## # ... with 50 more rows
```

Let's chain many of these verbs. What is the maximum GDP per continent across all years?

```
## Your turn
gapminder %>%
  group_by(continent, year) %>%
  summarize(max(gdpPercap))
```

```
## 'summarise()' regrouping output by 'continent' (override with '.groups' argument)
```

```
## # A tibble: 60 x 3
## # Groups:   continent [5]
##   continent year 'max(gdpPercap)'
##   <chr>      <dbl>      <dbl>
## 1 Africa     1952         4725.
## 2 Africa     1957         5487.
## 3 Africa     1962         6757.
## 4 Africa     1967        18773.
## 5 Africa     1972        21011.
## 6 Africa     1977        21951.
## 7 Africa     1982        17364.
## 8 Africa     1987        11864.
## 9 Africa     1992        13522.
## 10 Africa    1997        14723.
## # ... with 50 more rows
```

Find the maximum life expectancy for countries in Asia.

```
## Your turn
gapminder %>%
  filter(continent == "Asia") %>%
  group_by(continent, country) %>%
  summarize(max(lifeExp))
```

```
## 'summarise()' regrouping output by 'continent' (override with '.groups' argument)
```

```
## # A tibble: 33 x 3
## # Groups:   continent [1]
##   continent country      'max(lifeExp)'
##   <chr>      <chr>          <dbl>
## 1 Asia      Afghanistan      43.8
## 2 Asia      Bahrain          75.6
## 3 Asia      Bangladesh       64.1
## 4 Asia      Cambodia         59.7
```



```
## 5 Asia      China      73.0
## 6 Asia      Hong Kong China 82.2
## 7 Asia      India       64.7
## 8 Asia      Indonesia    70.6
## 9 Asia      Iran        71.0
## 10 Asia     Iraq        65.0
## # ... with 23 more rows
```

**Additional Data wrangling questions** Q. What other questions can you ask from the gapminder dataset (questions that you can answer by wrangling the data)? *Growth of life expectancy along with growth in life expectancy?*

**Data wrangling questions with another dataset** Q. PS1 was on nycflight data. Assuming that the problem set and data is fresh in your mind, what other questions can you ask from the flight dataset (questions that you can answer by wrangling the data)? And what dplyr functions would you use to answer that question? (*Example question: What is the typical delay of flights?*) *Are flights delayed more in the A.M. or P.M? Are there better times during the day? Functions to use are probably group\_by, max or min, and probably a filter*